Kevin Rajan

(331) 444-2203 | kevinsrajan@gmail.com | Chicago, IL | github.com/kvnloo | linkedin.com/in/kevinsrajan

Technical Skills

AI & LLMs: Hermes Agent, Claude Code, Codex, OpenCode, ACP, MCP, context engineering, evals
Languages: Python, C++, C, Java, Haskell
Data Analytics: pandas, Tableau, Cognos, PowerBI

Experience

LLM Training Specialist Jan. 2025 – Present
Outlier (Scale AI) *Remote*

- Designed and applied detailed rubrics for LLM evaluation, writing complex logic, math, and physics prompts to elicit reasoning failures and guide model improvement through process-supervision feedback. Performed red-teaming and fact-checking across STEM and general-knowledge domains, verifying claims with web sources and explaining weaknesses in model reasoning to refine future outputs.
- Conducted large-scale preference labeling and comparative ranking of multiple LLMs (RLHF, DPO, RLAIF) while drafting critiques that articulated why certain completions failed rubric standards. Developed multimodal and tool-use evaluations integrating text, image, and retrieval tasks to measure grounding, reliability, and reasoning depth across diverse model architectures.

Founder & CEO Jun. 2024 – Present
zero, LLC *Remote*

- Landed 41 commits on **NousResearch/hermes-agent** main (Aug.–Oct. 2026) across runtime correctness (cron, gateway watchdog, config signatures), context-compaction safety, kanban, desktop, and platform adapters; opened 194 PRs and 161 issues and commented on 1,272 threads upstream.
- Designed **z0**, a registry and installer for a personal agent stack (evidence-routed decisions, typed intent contracts, token economics, reproducible evals), and a verified open-source loop (leased claims, isolated worktrees, evidence receipts, human merge); 151 merged PRs in 44 external projects.
- Built **Hermes Tern**, a native terminal UI for Hermes (plan review, forked sessions, OpenUI previews, voice), and **Dash**, an Expo phone app plus Bun bridge over Tailscale that keeps one live thread per agent harness (Hermes, Codex, Claude Code, OMP, Grok); both in daily use.

Self-Directed Ultralearning Jul. 2023 – Jun. 2024
Sabbatical *Naperville, IL*

- Engineered a holistic performance optimization system to enhance cognitive and motor skills, applying an algorithmic feedback loop (Research, Implement, Analyze, Iterate) to systematically correlate and refine inputs, including sleep, nutrition, and specialized training protocols.
- Executed a long-term data-driven analysis project to benchmark and optimize rapid skill acquisition; successfully reverse-engineered complex, dynamic systems to identify and eliminate bottlenecks, culminating in validated performance metrics within the top 1–5% of a 25M+ user environment (Kovaaks, Valorant, Apex).

Senior Consultant, Technology Enablement Jun. 2021 – Jul. 2023
Slalom Consulting *Chicago, IL*

- Collaborated within a multi-disciplinary team to enhance client ($4B+ revenue finance firm) engagement while modernizing technology systems across 12 teams and verticals. Mitigated obstacles for business unit teams using enhanced software processes and workflows, promoting Agile/DevOps/CI/CD concepts by utilizing AWS, Jenkins, Terraform and Vault.
- Improved developer impact by creating tailored pipelines that publish biweekly golden Windows and RHEL releases, ultimately ensuring security standards within custom AMIs.

Projects

Quackles (Microduck product study) 2026

- Product page homage to the Hermes Agent poster: Blender Cycles plates in five themes and a 201 MP deep-zoom master served as a WebP tile pyramid.
- The same pyramid backs an infinite-canvas portfolio with semantic zoom: overview, summaries, details, then commit-level evidence.

Education

University of Illinois at Urbana-Champaign — Urbana-Champaign, IL
B.S. Engineering Physics, Minor in Computer Science — *Graduated May 2020*